

The AI4K12 Initiative: Developing National Guidelines for Teaching AI In K-12

David S. Touretzky
Carnegie Mellon University

Global SW Education Conference
Seoul, Korea. November 3-4, 2020.



#AI4K12

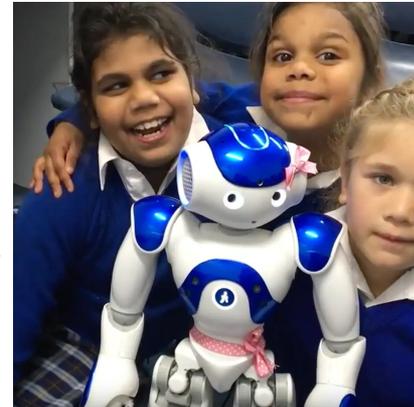


Why should children be learning about AI at all?

1. For the same reason they learn about electricity: *AI is powering the next industrial revolution.*

We need informed citizens who understand the issues raised by new technology, because AI is bringing huge economic and social changes.

2. Children are growing up with AI all around them. They should not regard it as magic.
3. Encourage thinking about AI-related careers.

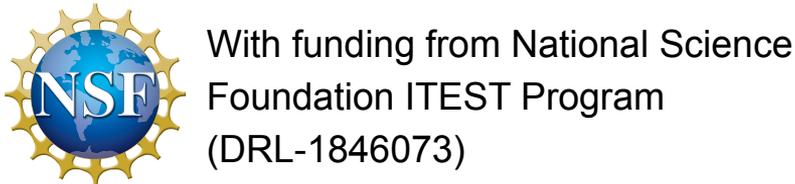


The AI4K12 Initiative, a joint project of:

AAAI (Association for the Advancement of Artificial Intelligence)



CSTA (Computer Science Teachers Association)



Carnegie Mellon University
School of Computer Science



AI4K12
.org

Initiative: Mission

- **Develop national guidelines for teaching AI in K-12**
 - Modeled after the CSTA standards for computing education.
 - Four grade bands: K-2, 3-5, 6-8, and 9-12
 - What should students know?
 - What should students be able to do?
- **Develop a curated AI resource directory for K-12 teachers**
- **Foster a K-12 AI Education Research and Practice Community & Resource Developers**



Steering Committee



Dave Touretzky
Carnegie Mellon
AI for K-12 Working Group
Chair



Christina Gardner-McCune
University of Florida
AI For K-12 Working
Group Co-Chair



Fred Martin
(emeritus member)
UMass Lowell
Former CSTA Chair of Board
of Directors



Deborah Seehorn
Co-Chair of CSTA
Standards Committee



K-12 Teacher Working Group Members



Grades K-2

Vicky Sedgwick (Lead)

Susan Amsler-Akacem

Dr. April DeGennaro

Melissa Unger (New)

Grades 3-5

Kelly Powers (Lead)

Dr. Marlo Barnett

Dr. Phillip Eaglin

Alexis Cobo (New)

Grades 6-8

Sheena Vaidyanathan (Lead)

Padmaja Bandaru

Josh Caldwell

Charlotte Dungan

Rachael Smith (New)

Grades 9-12

Jared Amalong (Lead)

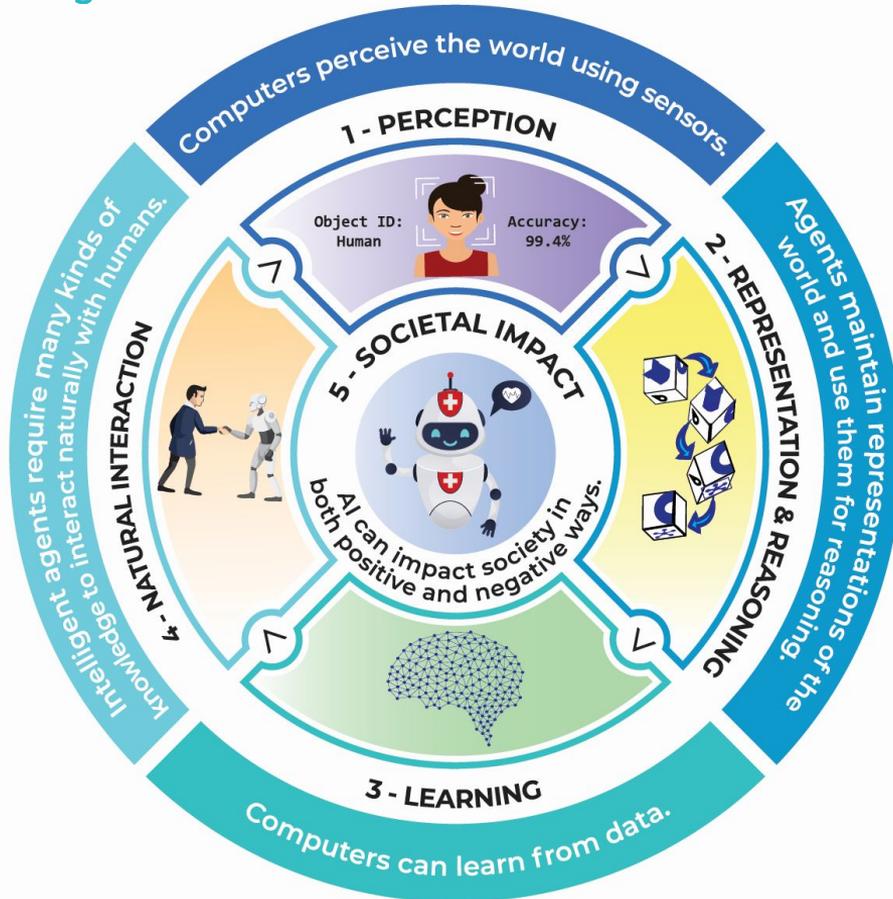
Dr. Smadar Bergman

Kate Lockwood

John Chapin (New)

Year 1 Alumni: Brian Stamford, Minsoo Park, Juan Palomares, Vincent Gregorio, Dianne O'Grady-Cunniff

Five Big Ideas in AI



Big Idea #2: Representation and Reasoning

Agents maintain representations of the world, and use them for reasoning.

- Types of representations
- Families of algorithms and the work they do
- Representation supports reasoning: algorithms operate on representations

Example Guidelines

- Create/design a representation of an (animal) classification system using a tree structure.
- Draw a search tree for tic-tac-toe
- Describe how AI representations support reasoning to answer questions
- Describe the differences between types of search algorithms



Big Idea #3: Learning

Computers can learn from data.

- Nature of learning
- Fundamentals of neural networks
- Data sets



Example Guidelines

- Modify an interactive machine learning project by training its model..
- Describe how algorithms and machine learning can exhibit biases.
- Identify bias in a training data set and extend the training set to address the bias
- Train a neural net (1-3 layers) using *TensorFlow Playground*
- Trace and experiment with a simple ML algorithm

Big Idea #4: Natural Interaction

Intelligent agents require many kinds of knowledge to interact naturally with humans.

- Natural language understanding
- Common sense reasoning
- Affective computing & interaction (e.g. with robots, or speech agents)
- Consciousness and philosophy of mind



Example Guidelines

- Recognize and label facial expressions into appropriate emotions (happiness, sadness, anger) and explain why they are labeled the way they are
- Experiment with software that recognizes emotions in facial expressions
- Construct a simple chatbot
- Describe some tasks where AI outperforms humans, and tasks where it does not
- Explain and give examples of how language can be ambiguous
- Reason about the nature of intelligence, and identify approaches to determining whether an agent is or is not intelligent.

Big Idea #5: Societal Impact (1 of 3)

“Artificial Intelligence can impact society in both positive and negative ways.”

- **Ethics of AI making decisions about people**

- Fairness and bias
- Transparency and explainability
- Accountability

Example Guidelines

- Critically explore the positive and negative impacts of an AI system.
- Describe ways that AI systems can be designed for inclusivity.



Machine Bias: ProPublica.org

Big Idea #5: Societal Impact (2 of 3)

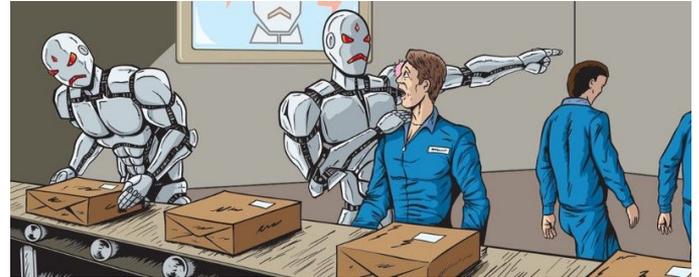
“Artificial Intelligence can impact society in both positive and negative ways.”

● Economic impacts of AI

- Increased productivity
- New types of services
- Reduction in of some types of jobs
- New career opportunities

Example Guidelines

- Design and explain how an AI system can be used to address a social issue.
- Understand tradeoffs in the design of AI systems and how decisions can have unintended consequences in the function of a system.



Big Idea #5: Societal Impact (3 of 3)

“Artificial Intelligence can impact society in both positive and negative ways.”

- **AI & Culture**

- Living with intelligent assistants and robot companions.
- Would you let your child travel unaccompanied in a self-driving car?
- New YouTube genre: self-driving car mishaps.

Example Guidelines

- Critically explore the positive and negative impacts of an AI system.
- Describe the debate about whether people should be polite to agents and robots.



Five Big Ideas in Artificial Intelligence

5. Societal Impact

AI can impact society in both positive and negative ways. AI technologies are changing the ways we work, travel, communicate, and care for each other. But we must be mindful of the harms that can potentially occur. For example, biases in the data used to train an AI system could lead to some people being less well served than others. Thus, it is important to discuss the impacts that AI is having on our society and develop criteria for the ethical design and deployment of AI-based systems.

4. Natural Interaction

Intelligent agents require many kinds of knowledge to interact naturally with humans. Agents must be able to converse in human languages, recognize facial expressions and emotions, and draw upon knowledge of culture and social conventions to infer intentions from observed behavior. All of these are difficult problems. Today's AI systems can use language to a limited extent, but lack the general reasoning and conversational capabilities of even a child.

1. Perception

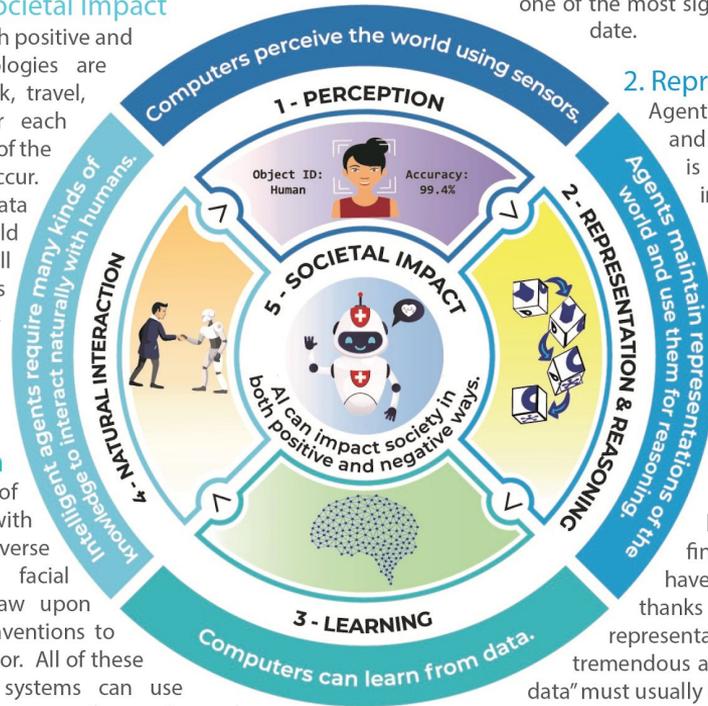
Computers perceive the world using sensors. Perception is the process of extracting meaning from sensory signals. Making computers "see" and "hear" well enough for practical use is one of the most significant achievements of AI to date.

2. Representation & Reasoning

Agents maintain representations of the world and use them for reasoning. Representation is one of the fundamental problems of intelligence, both natural and artificial. Computers construct representations using data structures, and these representations support reasoning algorithms that derive new information from what is already known. While AI agents can reason about very complex problems, they do not think the way a human does.

3. Learning

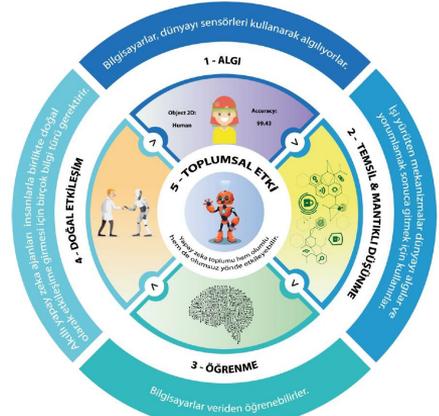
Computers can learn from data. Machine learning is a kind of statistical inference that finds patterns in data. Many areas of AI have progressed significantly in recent years thanks to learning algorithms that create new representations. For the approach to succeed, tremendous amounts of data are required. This "training data" must usually be supplied by people, but is sometimes acquired by the machine itself.



Adoption of the Big Ideas

- Now being adopted by curriculum developers in the US and elsewhere.
- Translations available in Chinese, Korean, Portuguese, Hebrew, and Turkish. Spanish, Japanese coming soon.

Turkish



Chinese

人工智能的五大理念

5. 社会影响
AI的应用对社会既有正面影响也有负面影响。人工智能技术正在改变我们工作、出行、沟通、和相互照顾的方式，但我们必须注意其所能带来的危害。例如，若用于训练人工智能系统的数据存在偏见，可能会导致部分人受到的服务质量低于其他人。因此，讨论AI对我们社会的影响，并根据相关系统在道德层面的设计以及应用来制定标准是重要的。

4. 人机交互
智能代理需要多种知识才能与人类自然交互。为了与人类自然地交互，智能代理必须能够用人类语言交谈，识别面部表情和情绪，并利用文化和人类社会的知识来推断所观察到的人类行为的意图。所有这些问题想要解决都不容易。今天的人工智能系统可以在有限的程度上使用语言，但其综合推理和会话能力却不如一般的人类儿童。



2. 表示与推理
智能代理 (能够) 保持对现实世界的表示,并用他们进行推理。表示是自然智能和人工智能的基本问题之一。计算机使用数据结构来构建表示。这些表示辅助推理算法。这些推理算法从而从已知信息中推导出新的信息。虽然智能代理可以推理非常复杂的问题,但他们并不像人类一样思考问题。

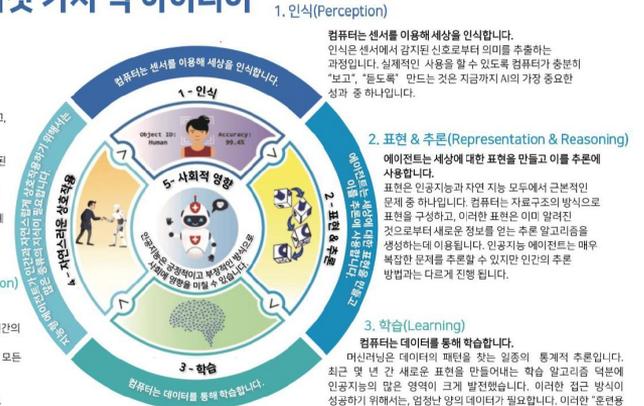
3. 机器学习
计算机可以从数据中学习。机器学习是一种在数据中找到规律的统计推断。近年来,由于一些学习算法创造了新的表示,AI的许多领域都取得了显著进展。这种方法的成功需要大量的数据。这些“训练数据”通常必须由人们提供,但有时也可以由机器自身获取。

Korean

인공지능에 관한 다섯 가지 빅 아이디어

5. 사회적 영향(Social Impact)
인공지능은 긍정적이고 부정적인 방식으로 사회에 영향을 미칠 수 있습니다. 인공지능 기술은 우리가 일하고, 여행하고, 의사 소통하고, 서로를 돕는 방식을 변화시키고 있습니다. 그러나 우리는 잠재적으로 발생할 수 있는 위험에 유의해야 합니다. 예를 들어, 인공지능 시스템을 훈련하는 데 편향된 데이터를 이용하면, 인공지능은 다른 사람들과 다르게 된 지평을 발견하는 경우가 발생할 수 있습니다. 그러므로, 인공지능이 우리 사회에 미치는 영향에 대해 논의할 필요가 있고 인공지능 기반 시스템의 윤리적 설계 및 배치에 관한 기술을 개발하는 것이 중요합니다.

4. 자연스러운 상호작용(Natural Interaction)
지능형 에이전트가 인간과 자연스럽게 상호작용하기 위해서는 많은 종류의 지식이 필요합니다. 에이전트가 관찰된 행동의 의도를 추론하기 위해서는 인간의 언어로 대화하고, 얼굴 표정과 감정을 인식하며, 사회적 관습과 문맥에 대한 지식을 활용할 수 있어야 합니다. 이 모든 것들은 매우 어려운 문제들입니다. 오늘날의 인공지능 시스템은 제한된 범위에서 언어를 사용할 수 있지만, 일반적인 주제를 대화 능력은 아이보다도 부족합니다.



1. 인식(Perception)

컴퓨터는 센서를 이용해 세상을 인식합니다. 인식은 센서에서 감지된 신호로부터 의미를 추출하는 과정입니다. 실제적인 사용을 할 수 있도록 컴퓨터가 충분히 알고, "틀들" 만드는 것은 지금까지의 가장 중요한 성과 중 하나입니다.

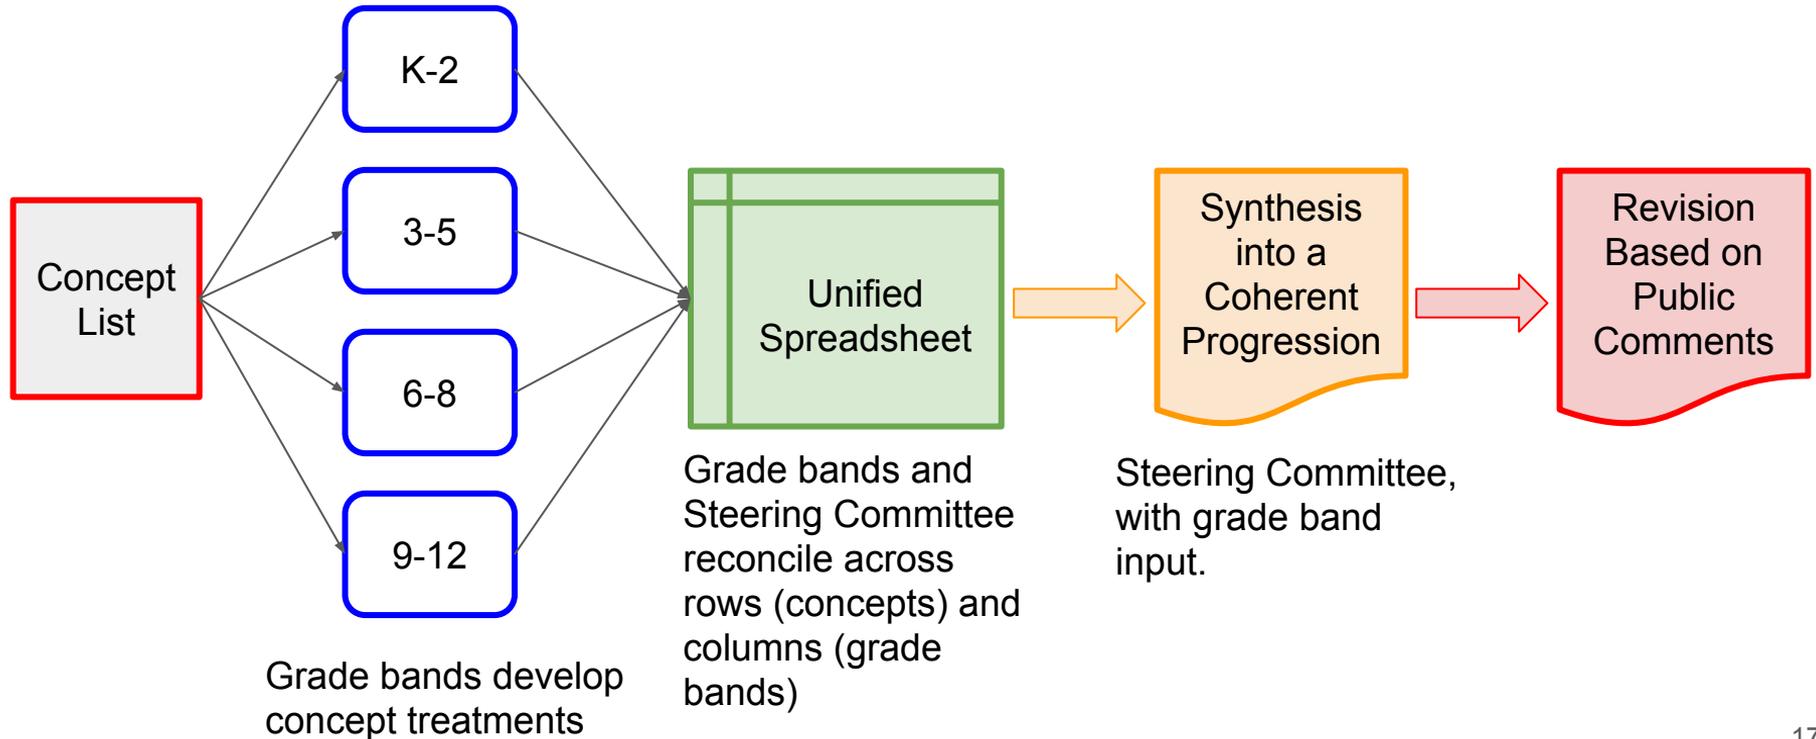
2. 표현 & 추론(Representation & Reasoning)

에이전트는 세상에 대한 표현을 만들고 이를 추론에 사용합니다. 표현은 인공지능과 지능 모두에서 근본적인 문제 중 하나입니다. 컴퓨터는 자료구조의 방식으로 표현을 구성하고, 이러한 표현은 이미 알려진 것으로부터 새로운 정보를 얻는 추론 알고리즘을 생성하는데 이용됩니다. 인공지능 에이전트는 매우 복잡한 문제를 추론할 수 있지만 인간의 추론 방법과는 다르게 진행합니다.

3. 학습(Learning)

컴퓨터는 데이터를 통해 학습합니다. 머신러닝은 데이터의 패턴을 찾는 일종의 통계 추론입니다. 최근 몇 년 간 새로운 표현을 만들어내는 학습 알고리즘 덕분에 인공지능의 많은 영역이 크게 발전했습니다. 이러한 접근 방식이 성공하기 위해서는, 엄청난 양의 데이터가 필요합니다. 이러한 "훈련 데이터(training data)"는 일반적으로 사람이 제공해야 하지만, 때로는 기계 스스로 수집하기도 합니다.

Drafting the Guidelines for One Big Idea



Big Idea #1: Perception	Computers perceive the world using sensors.		Perception is the extraction of meaning from sensory information using knowledge.	The transformation from signal to meaning takes place in stages, with increasingly abstract features and higher level knowledge applied at each stage.	LO = Learning Objective: what students should be able to do. EU = Enduring Understanding: what students should know.
Concept	K-2	3-5	6-8	9-12	
Sensing (Living Things) 1-A-i	LO: Identify human senses and sensory organs. EU: People experience the world through sight, hearing, touch, taste, and smell.	LO: Compare human and animal perception. EU: Some animals experience the world differently than people do. Unpacked: Bats and dolphins use sonar. Bees can see ultraviolet. Rats have no color vision; dogs are red-green colorblind. Dogs and rats can hear higher frequencies than humans.	LO: Give examples of how humans combine information from multiple modalities. EU: People can exploit correlations between senses, such as sight and sound, to make sense of ambiguous signals. Unpacked: In a noisy environment, speech is more understandable when the speaker's mouth is visible. People learn the sounds associated with various actions (such as dropping an object) and can recognize when the sound doesn't match their expectation.	N/A – for AI purposes, this topic has already been adequately addressed in the lower grade bands. Other courses, such as biology or an elective on sensory psychology, could go into more detail about topics such as taste, smell, proprioception, and vestibular organs. Possible enrichment material: look at optical illusions (Müller-Lyer illusion, Kanizsa triangle) and ask which ones are computer vision systems also subject to.	
Sensing (Computer Sensors) 1-A-ii	LO: Locate and identify sensors (camera, microphone) on computers, phones, robots, and other devices. EU: Computers "see" through video cameras and "hear" through microphones.	LO: Illustrate how computer sensing differs from human sensing. EU: Most computers have no sense of taste, smell, or touch, but they can sense some things that humans can't, such as infrared emissions, extremely low or high frequency sounds, or magnetism.	LO: Give examples of how intelligent agents combine information from multiple sensors. EU: Self driving cars combine computer vision with radar or lidar imaging. GPS measurement, and accelerometer data to form a detailed representation of the environment and their motion through it.	LO: Describe the limitations and advantages of various types of computer sensors. EU: Sensors are devices that measure physical phenomena such as light, sound, temperature, or pressure. Unpacked: Cameras have limited resolution, dynamic range, and spectral sensitivity. Microphones have limited sensitivity and frequency response. Signals may be degraded by noise, such as a microphone in a noisy environment. Some sensors can detect things that people cannot, such as infrared or ultraviolet imagery, or ultrasonic sounds.	
Sensing (Digital Encoding) 1-A-iii	N/A	LO: Explain how images are represented digitally in a computer. EU: Images are encoded as 2D arrays of pixels, where each pixel is a number indicating the brightness of that piece of the image, or an RGB value indicating the brightness of the red, green, and blue components of that piece.	LO: Explain how sounds are represented digitally in a computer. EU: Sounds are digitally encoded by sampling the waveform at discrete points (typically several thousand samples per second), yielding a series of numbers.	LO: Explain how radar, lidar, GPS, and accelerometer data are represented. EU: Radar and lidar do depth imaging: each pixel is a depth value. GPS triangulates position using satellite signals and gives a location as longitude and latitude. Accelerometers measure acceleration in 3 orthogonal dimensions. Unpacked: Radar and lidar measure distance as the time for a reflected signal to return to the transceiver. GPS determines position by triangulating precisely timed signals from three or more satellites. Accelerometers use orthogonally oriented strain gauges to measure acceleration in three dimensions.	

Page 1 (of 4)
of the draft
guidelines for
Big Idea #1:
Perception

Big Idea #1: Perception

Computers perceive the world using sensors.

Perception is the extraction of meaning from sensory information using knowledge.

The transformation from signal to meaning takes place in stages, with increasingly abstract features and higher level knowledge applied at each stage.

AIHK12
Draft Big Idea 1 - Progression Chart
www.AIHK12.org

Big Idea #1: Perception
Cognitive processes that extract meaning from sensory information using knowledge.

Grade	K-2	3-5	6-8	9-12
Meaning	KS-10-12-1-1-1: Identify and describe the basic elements of a scene (e.g., objects, colors, shapes, locations, movements) using the senses.	3S-10-12-1-1-1: Identify and describe the basic elements of a scene (e.g., objects, colors, shapes, locations, movements) using the senses.	6S-10-12-1-1-1: Identify and describe the basic elements of a scene (e.g., objects, colors, shapes, locations, movements) using the senses.	9S-10-12-1-1-1: Identify and describe the basic elements of a scene (e.g., objects, colors, shapes, locations, movements) using the senses.
Knowledge	KS-10-12-1-1-2: Use basic knowledge to identify and describe the basic elements of a scene (e.g., objects, colors, shapes, locations, movements) using the senses.	3S-10-12-1-1-2: Use basic knowledge to identify and describe the basic elements of a scene (e.g., objects, colors, shapes, locations, movements) using the senses.	6S-10-12-1-1-2: Use basic knowledge to identify and describe the basic elements of a scene (e.g., objects, colors, shapes, locations, movements) using the senses.	9S-10-12-1-1-2: Use basic knowledge to identify and describe the basic elements of a scene (e.g., objects, colors, shapes, locations, movements) using the senses.
Skills	KS-10-12-1-1-3: Use basic skills to identify and describe the basic elements of a scene (e.g., objects, colors, shapes, locations, movements) using the senses.	3S-10-12-1-1-3: Use basic skills to identify and describe the basic elements of a scene (e.g., objects, colors, shapes, locations, movements) using the senses.	6S-10-12-1-1-3: Use basic skills to identify and describe the basic elements of a scene (e.g., objects, colors, shapes, locations, movements) using the senses.	9S-10-12-1-1-3: Use basic skills to identify and describe the basic elements of a scene (e.g., objects, colors, shapes, locations, movements) using the senses.

V3.1 - Released May 28, 2020
Subject to change based on public feedback

1-A-i: Sensing in Living Things

K-2

LO: Identify human senses and sensory organs.

EU: People experience the world through sight, hearing, touch, taste, and smell.

3-5

LO: Compare human and animal perception.

EU: Some animals experience the world differently than people do.

Unpacked: Bats and dolphins use sonar. Bees can see ultraviolet. Rats have no color vision...

LO (Learning Objective): What students should be able to do.

EU (Enduring Understanding): What students should know.

Big Idea #3: Learning

Computers can learn from data.

Machine learning allows a computer to acquire behaviors without people explicitly programming those behaviors.

Learning of new behaviors is brought about by changes in the internal representations of a reasoning model, such as a neural network or decision tree.

Large amounts of training data are required to narrow down the learning algorithm's choices when the reasoning model is capable of a great variety of behaviors.

The reasoning model constructed by the machine learning algorithm can be applied to new data to solve problems or make decisions.

Big Idea #3 Concept List

3-A: Nature of Learning

- 3-A-i: Humans vs. Machines
- 3-A-ii: Finding Patterns in Data
- 3-A-iii: Training a Model
- 3-A-iv: Constructing a Reasoner
- 3-A-v: Adjusting Parameters
- 3-A-vi: Learning from Experience

3-B: Neural Networks

- 3-B-i: Structure of a Neural Network
- 3-B-ii: Weight Adjustment

3-C: Datasets

- 3-C-i: Feature Sets
- 3-C-ii: Large Datasets
- 3-C-iii: Bias

What's Next?

- October 2020: First release of Big Idea #3 progression chart for public comment.
- November 2020: Revised Big Idea #1 progression chart.
- Early 2021: drafts of progression charts for Big Ideas #4, 5, and 2.
- February 2021: EAAI-21 will have a special track: “Demos, Software Tools, and Activities for Teaching AI in K-12”.

Join Us in Developing the Guidelines, or Help Grow the Community of AI Resource Developers

Visit us:

<http://AI4K12.org>

Join the mailing list:

<https://aaai.org/Organization/mailing-lists.php>

